

# Stanford CME 241 (Winter 2021) - Assignment 3

## Frog on Lilypad (code part)

Chih-Hsuan 'Carolyn' Kao (chkao831@stanford.edu)

Feb 28th, 2021

In [1]:

```
import sys
sys.path.append('/Users/chih-hsuankao/Desktop/CME241/RL-book/')

from rl.distribution import Categorical, Constant
from rl.dynamic_programming import (
    evaluate_mrp_result,
    policy_iteration_result,
    value_iteration_result
)
from rl.markov_decision_process import (
    FiniteMarkovDecisionProcess,
    FinitePolicy,
    StateActionMapping,
)
from rl.markov_process import (
    Transition,
    RewardTransition,
    FiniteMarkovProcess,
    Optional,
    FiniteMarkovRewardProcess,
)
```

```
/Users/chih-hsuankao/.pyenv/versions/anaconda3-2019.03/lib/python3.7/site-packages/scipy/__init__.py:137: UserWarning: NumPy 1.16.5 or above is required for this version of SciPy (detected version 1.16.2)
  UserWarning)
```

In [2]:

```
from dataclasses import dataclass
import itertools
import matplotlib.pyplot as plt
from typing import Mapping, Dict, Tuple, List
```

Consider an array of  $n + 1$  lily pads on a pond, numbered  $0$  to  $n$ . A frog sits on a lily pad other than the lily pads numbered  $0$  or  $n$ . When on lily pad  $i$  ( $1 \leq i \leq n - 1$ ), the frog can croak one of two sounds A or B.

If it croaks A when on lily pad  $i$  ( $1 \leq i \leq n - 1$ ), it is thrown to lily pad  $i - 1$  with probability  $\frac{i}{n}$  and is thrown to lily pad  $i + 1$  with probability  $\frac{n - i}{n}$ . If it croaks B when on

lily pad  $i$  ( $1 \leq i \leq n - 1$ ), it is thrown to one of the lily pads  $0, \dots, i-1, i+1, \dots, n$  with uniform probability  $\frac{1}{n}$ . A snake, perched on lily pad  $0$ , will eat the frog if the frog lands on lily pad  $0$ . The frog can escape the pond (and hence, escape the snake!) if it lands on lily pad  $n$ .

What should the frog croak when on each of the lily pads  $1, 2, \dots, n - 1$ , in order to maximize the probability of escaping the pond (i.e., reaching lily pad  $n$  before reaching lily pad  $0$ )? Although there are more than one ways of solving this problem, we'd like to solve it by modeling it as an MDP and identifying the Optimal Policy.

Write code to model this MDP as an instance of the `FiniteMarkovDecisionProcess` class. We have learnt that there exists an optimal deterministic policy, and there are  $2^{n-1}$  possible deterministic policies for this problem. Write code to create each of these  $2^{n-1}$  deterministic policies (as instances of `FinitePolicy` class), create a policy-implied Finite MRP for each of these deterministic policies (using the `apply finite policy` method of `FiniteMarkovDecisionProcess` class), and evaluate the Value Function for each of those implied Finite MRPs. This should give you the Optimal Value Function and the Optimal Deterministic Policy.

In [3]:

```
@dataclass(frozen=True)
class FrogState:
    position: int
```

In [4]:

```
FrogJumpMap = StateActionMapping[FrogState, int]
```

In [5]:

```
class FrogMDP(FiniteMarkovDecisionProcess[FrogState, str]):

    def __init__(
        self,
        num_pad: int = 10,
    ):
        self.num_pad = num_pad

        super().__init__(self.get_action_transition_reward_map())

    def get_action_transition_reward_map(self) -> StateActionMapping[FrogState, str]:

        d: Dict[FrogState, Dict[str, Categorical[Tuple[FrogState, float]]]] = {}

        # ref: https://github.com/coverdrive/MDP-DP-RL/blob/master/src/examples/exam_problems/frog_lilypad.py
        for i in range(1, self.num_pad):

            d1: Dict[str, Categorical[Tuple[FrogState, float]]] = {}

            # Croak A
            d1["A"] = Categorical({(FrogState(i - 1), 0.):
                                   i / self.num_pad,
                                   (FrogState(i + 1), 1. if i == self.num_
pad-1 else 0.):
```

```

                                (self.num_pad - i) /self.num_pad})
        # Croak B
        d1["B"] = Categorical({(FrogState(j), 1. if j == self.num_pad
else 0.):
                                1/self.num_pad for j in range(self.
num_pad + 1) if j != i})

        d[FrogState(i)] = d1

        d[FrogState(self.num_pad)] = None
        d[FrogState(0)] = None

        return d

    def rewardf(
        self,
        current_pad: int,
        num_pad: int
    ):
        if current_pad == num_pad:
            return 1.

        elif current_pad == 0:
            return -1.

        else:
            return 0.

```

In [6]:

```

if __name__ == '__main__':

    gamma = 0.8
    pad = 10

    si_mdp: FiniteMarkovDecisionProcess[FrogState, int] =\
        FrogMDP(
            num_pad = pad
        )

    print("MDP Transition Map")
    print("-----")
    print(si_mdp)

    policies = list(itertools.product([0, 1], repeat = pad - 1))
    #print(policies)

    # For each deterministic policy
    for policy in policies:
        # print("A Deterministic Policy:")
        fdp: FinitePolicy[FrogState, int] =\
            FinitePolicy(
                {FrogState(padnum):
                 Constant(policy[padnum - 1]) for padnum in range(1, pa
d)}
            )

        # commented out to avoid long output; uncomment the line below as
        needed
        # print(fdp)

```

```

print("Optimal Value Function and Optimal Policy")
print("-----")

opt_vf_vi, opt_policy_vi = value_iteration_result(si_mdp, gamma=gamma)
print(opt_vf_vi)
print(opt_policy_vi)

```

#### MDP Transition Map

-----

From State FrogState(position=1):

With Action A:

To [State FrogState(position=0) and Reward 0.000] with Probability 0.100

To [State FrogState(position=2) and Reward 0.000] with Probability 0.900

With Action B:

To [State FrogState(position=0) and Reward 0.000] with Probability 0.100

To [State FrogState(position=2) and Reward 0.000] with Probability 0.100

To [State FrogState(position=3) and Reward 0.000] with Probability 0.100

To [State FrogState(position=4) and Reward 0.000] with Probability 0.100

To [State FrogState(position=5) and Reward 0.000] with Probability 0.100

To [State FrogState(position=6) and Reward 0.000] with Probability 0.100

To [State FrogState(position=7) and Reward 0.000] with Probability 0.100

To [State FrogState(position=8) and Reward 0.000] with Probability 0.100

To [State FrogState(position=9) and Reward 0.000] with Probability 0.100

To [State FrogState(position=10) and Reward 1.000] with Probability 0.100

From State FrogState(position=2):

With Action A:

To [State FrogState(position=1) and Reward 0.000] with Probability 0.200

To [State FrogState(position=3) and Reward 0.000] with Probability 0.800

With Action B:

To [State FrogState(position=0) and Reward 0.000] with Probability 0.100

To [State FrogState(position=1) and Reward 0.000] with Probability 0.100

To [State FrogState(position=3) and Reward 0.000] with Probability 0.100

To [State FrogState(position=4) and Reward 0.000] with Probability 0.100

To [State FrogState(position=5) and Reward 0.000] with Probability 0.100

To [State FrogState(position=6) and Reward 0.000] with Probability 0.100

To [State FrogState(position=7) and Reward 0.000] with Probability 0.100

To [State FrogState(position=8) and Reward 0.000] with Probability 0.100

probability 0.100

To [State FrogState(position=9) and Reward 0.000] with Probability 0.100

To [State FrogState(position=10) and Reward 1.000] with Probability 0.100

From State FrogState(position=3):

With Action A:

To [State FrogState(position=2) and Reward 0.000] with Probability 0.300

To [State FrogState(position=4) and Reward 0.000] with Probability 0.700

With Action B:

To [State FrogState(position=0) and Reward 0.000] with Probability 0.100

To [State FrogState(position=1) and Reward 0.000] with Probability 0.100

To [State FrogState(position=2) and Reward 0.000] with Probability 0.100

To [State FrogState(position=4) and Reward 0.000] with Probability 0.100

To [State FrogState(position=5) and Reward 0.000] with Probability 0.100

To [State FrogState(position=6) and Reward 0.000] with Probability 0.100

To [State FrogState(position=7) and Reward 0.000] with Probability 0.100

To [State FrogState(position=8) and Reward 0.000] with Probability 0.100

To [State FrogState(position=9) and Reward 0.000] with Probability 0.100

To [State FrogState(position=10) and Reward 1.000] with Probability 0.100

From State FrogState(position=4):

With Action A:

To [State FrogState(position=3) and Reward 0.000] with Probability 0.400

To [State FrogState(position=5) and Reward 0.000] with Probability 0.600

With Action B:

To [State FrogState(position=0) and Reward 0.000] with Probability 0.100

To [State FrogState(position=1) and Reward 0.000] with Probability 0.100

To [State FrogState(position=2) and Reward 0.000] with Probability 0.100

To [State FrogState(position=3) and Reward 0.000] with Probability 0.100

To [State FrogState(position=5) and Reward 0.000] with Probability 0.100

To [State FrogState(position=6) and Reward 0.000] with Probability 0.100

To [State FrogState(position=7) and Reward 0.000] with Probability 0.100

To [State FrogState(position=8) and Reward 0.000] with Probability 0.100

To [State FrogState(position=9) and Reward 0.000] with Probability 0.100

To [State FrogState(position=10) and Reward 1.000] with Probability 0.100

From State FrogState(position=5):

With Action A:





To [State FrogState(position=3) and Reward 0.000] with Probability 0.100  
 To [State FrogState(position=4) and Reward 0.000] with Probability 0.100  
 To [State FrogState(position=5) and Reward 0.000] with Probability 0.100  
 To [State FrogState(position=6) and Reward 0.000] with Probability 0.100  
 To [State FrogState(position=7) and Reward 0.000] with Probability 0.100  
 To [State FrogState(position=8) and Reward 0.000] with Probability 0.100  
 To [State FrogState(position=10) and Reward 1.000] with Probability 0.100  
 FrogState(position=10) is a Terminal State  
 FrogState(position=0) is a Terminal State

#### Optimal Value Function and Optimal Policy

-----  
 {FrogState(position=1): 0.2824061058293976, FrogState(position=2): 0.2824061058293976, FrogState(position=3): 0.2824061058293976, FrogState(position=4): 0.2824061058293976, FrogState(position=5): 0.2824061058293976, FrogState(position=6): 0.2824061058293976, FrogState(position=7): 0.2824061058293976, FrogState(position=8): 0.2824061058293976, FrogState(position=9): 0.30332433866457054}  
 For State FrogState(position=1):  
   Do Action B with Probability 1.000  
 For State FrogState(position=2):  
   Do Action B with Probability 1.000  
 For State FrogState(position=3):  
   Do Action B with Probability 1.000  
 For State FrogState(position=4):  
   Do Action B with Probability 1.000  
 For State FrogState(position=5):  
   Do Action B with Probability 1.000  
 For State FrogState(position=6):  
   Do Action B with Probability 1.000  
 For State FrogState(position=7):  
   Do Action B with Probability 1.000  
 For State FrogState(position=8):  
   Do Action B with Probability 1.000  
 For State FrogState(position=9):  
   Do Action A with Probability 1.000