

## HW11

### 1.

Here, we consider a simple case of Monte-Carlo Prediction where the MRP consists of a finite state space with the non-terminal states  $\mathcal{N} = \{s_1, s_2, \dots, s_m\}$ . In this case, we represent the Value Function of the MRP in dictionary of (state, expected return) pairs, the relevant code is sketched out as the following,

```
S = TypeVar('S')

def tb_mc_pred(traces: Iterable[Iterable[mp.TransitionStep[S]]]
               gamma: float,
               tol: float=1e-08,
               vf: Dict[S,float])
    ->List[Mapping[S,float]]:

    episodes=[returns(trace,gamma,tol) for trace in traces]

    vf:List[Mapping[S,float]]=[]
    count:Mapping[S,int]=defaultdict(lambda:0)

    for episode in episodes:
        if(len(vf) == 0):
            vn: Mapping[S,float] = defaultdict(lambda:0)
        else:
            vn: Mapping[S,float] = {key:val for (key,val) in vf[-1].items()}
        for epi in episode:
            count[epi.state]+=1
            vn[epi.state] = vn[epi.state] + 1.0/(count[epi.state])*(epi.return_-vn[epi.state])
        vf.append(vn)
    return vf
```

In the code, we make update to the value function for state  $s_i$  by the weight for its latest trace experience return times the difference between the latest trace experience return and the current Value Function estimate. The formula is mathematically presented by:

$$V_n(s_i) = \frac{n-1}{n} \cdot V_{n-1}(s_i) + \frac{1}{n} \cdot Y_i^{(n)} = V_{n-1}(s_i) + \frac{1}{n} \cdot \left( Y_i^{(n)} - V_{n-1}(s_i) \right)$$